

TUSHAR LIMBASIYA

Ahmedabad, Gujarat, India

+91 6354197829

tusharlimbasiya200@gmail.com

linkedin.com/in/tushar-limbasiya-359917214

Senior Software Engineer | Python · Full-Stack · AI/ML · Cloud

PROFESSIONAL SUMMARY

Senior Software Engineer with 6+ years of full-stack commercial experience delivering high-impact Python applications across Finance, Healthcare, Enterprise Data Intelligence, and SaaS domains. Expert in advanced Python scripting, backend architecture (Django, FastAPI, Flask), and frontend development (React, JavaScript) with a proven track record of designing and deploying scalable, production-grade systems on AWS, Azure, and GCP. Deep expertise in Generative AI, LLM API integration, RAG pipelines, and Prompt Engineering — productionizing AI workflows for real-world commercial applications. Highly skilled in end-to-end deployment: Docker, Kubernetes, Terraform, CI/CD pipelines, and cloud infrastructure management. Demonstrated ability to lead technical solutions independently for distributed remote teams across international environments, owning the full stack from database design and RESTful API development through frontend delivery and DevOps automation.

PROFESSIONAL EXPERIENCE

Software Developer | Infoanalytica

January 2024 – January 2026 Ahmedabad, Gujarat, India

- Architected and built backend systems using Django REST Framework to power document intelligence and VLM-based workflows, enabling seamless data exchange across distributed microservices and supporting enterprise-grade LLM-powered reporting at scale.
- Owned infrastructure end-to-end — designed scalable database architectures, deployed containerized microservices on AWS EC2 using Docker and Kubernetes, and established monitoring and reliability practices ensuring 99.9% uptime for production ML pipelines.
- Designed and operated high-throughput, low-latency data processing pipelines using Python, Celery, and Kafka, establishing best practices for CI/CD, observability, and debugging across 4+ distributed production platforms.
- Managed distributed systems including PostgreSQL and MongoDB with parallel processing, Redis caching, and optimized indexing strategies — supporting high-scale analytics, ML model training, and enterprise production workloads.
- Integrated RESTful APIs and third-party services to synchronize real-time business events at enterprise-grade production scale, implementing performance optimization techniques that reduced manual reconciliation overhead by 40%.
- Architected a cloud-native, multi-platform orchestration system using microservices architecture and adapter patterns, enabling centralized analytics and productionizing ML/VLM workflows across four distributed production environments.
- Improved production reliability by 30% through ownership of incident response and performance optimization — designed distributed task queuing with retry logic, deployed monitoring/alerting infrastructure, and established debugging best practices using observability tools.

Software Engineer | confidential

January 2020 – Present Ahmedabad, Gujarat, India

- Engineered advanced Python scripting solutions and automation frameworks for remote-first international clients across Finance and Healthcare verticals, independently owning the full software development lifecycle from architecture through production deployment.
- Designed and deployed full-stack applications — Python/Django/FastAPI backends with React and JavaScript frontends — delivering end-to-end features across financial data platforms and healthcare workflow management systems for globally distributed teams.
- Led end-to-end deployment pipelines using Docker, Kubernetes, and Terraform on AWS and Azure, establishing CI/CD best practices with GitHub Actions and Jenkins — ensuring zero-downtime releases and 99.9% uptime SLAs for production systems serving international users.
- Integrated Generative AI, LLM APIs, and ML model inference into production backend services, building intelligent automation workflows using Python scripting for data extraction, transformation, and analytics delivery — directly supporting Data Intelligence objectives for remote stakeholders.

EDUCATION

Bachelor of Technology in Computer Engineering

Silver Oak University, Gujarat, India | August 2020 – May 2024 | GPA: 9.7 / 10

SKILLS

Programming & Languages	Python (Expert — 6+ yrs), JavaScript (ES6+), TypeScript, SQL, Shell Scripting, Bash, .NET (Familiar)
Frameworks & APIs	Django, Django REST Framework, FastAPI, Flask, React.js, AJAX, RESTful APIs, GraphQL, LLM API Integration, Generative AI Frameworks
Cloud & DevOps	AWS (EC2, Lambda, S3, RDS, CloudWatch, IAM), GCP, Azure DevOps, Docker, Kubernetes, Helm, Terraform, Jenkins, GitHub Actions, CI/CD Pipelines, GitHub Actions, CI/CD Pipelines
Databases	PostgreSQL, MySQL, MongoDB, ClickHouse, Redis, SQLite, Database Design, Query Optimization, Data Modeling, Server-Side Architecture

AI / ML & Data	Pandas, NumPy, Scikit-learn, Generative AI, RAG Pipelines, Prompt Engineering, LLM/VLM Integration, ML Model Deployment, Data Preprocessing, Feature Engineering, ML Inference Pipelines, LLM/VLM Integration, Document Intelligence, Generative AI, RAG Pipelines
Tools & Practices	Git, GitHub, Jira, Confluence, Swagger/OpenAPI, Postman, Redis, Celery, Kafka, Web Scraping (Scrapy, BeautifulSoup, Selenium), FHIR/HL7 Awareness, FinTech API Integration, Debugging, Agile/Scrum, Microservices Architecture, High-Throughput Service Design
Soft Skills	Technical Leadership, Cross-functional Collaboration, Agile/Scrum, Remote Team Management, Client Communication, Problem Solving, Mentoring, Documentation

PROJECTS

Dealer Call Automation & Analytics Platform

- ▶ Architected and built a high-scale, low-latency call automation backend using Python and Django REST Framework, operating across multiple distributed platforms with cloud-native deployment supporting real-time analytics and enterprise production workloads.
- ▶ Designed and operated high-throughput services with distributed queues (Celery, Kafka) to ingest and normalize call event logs at scale, storing enriched datasets in PostgreSQL for ML/VLM workflow consumption in production environments.
- ▶ Owned backend software stack decisions and contributed to Solution Roadmap planning, architecting server-side application logic with RESTful API design patterns supporting real-time analytics and performance monitoring.
- ▶ Owned infrastructure reliability and performance optimization — implemented modular microservices architecture with automated task scheduling, monitoring/alerting, and proactive debugging to ensure enterprise-grade system resilience.

MiraDoor – Property Management Platform

- ▶ Architected a comprehensive backend platform using Django with three-tiered microservices architecture, deployed across cloud and on-prem environments using Docker and Kubernetes, ensuring scalability and production-grade separation for ML/VLM workflows.
- ▶ Maintained deep knowledge of database systems — designed optimized PostgreSQL schemas with data storage strategies, constraints, and indexing for high-performance, logic-heavy backend systems in a commercial project context.
- ▶ Solved technical issues autonomously using pragmatic approach to problem-solving — implemented software design patterns with strong troubleshooting skills, ensuring maintainability, modularity, and comprehensive unit test coverage.
- ▶ Designed and deployed versioned RESTful APIs using Django REST Framework with Infrastructure-as-Code practices (Terraform), comprehensive Swagger documentation, and CI/CD pipelines supporting data ingestion and ML model inference at scale.

InfoDialer – Predictive Dialing Platform

- ▶ Architected and built InfoDialer backend from R&D prototypes to enterprise-ready production system on AWS EC2, supporting 1,000+ concurrent users with high-throughput, low-latency call routing and distributed network rules at scale.
- ▶ Developed a secure Admin Portal with Role-Based Access Control (RBAC), applying cybersecurity training and security best practices for user management, access auditing, and backend security architecture in a commercial production environment.
- ▶ Proficient in API integrations — integrated third-party REST APIs end-to-end with robust error handling, maintaining deep knowledge of API integration patterns and server-side architecture for seamless backend service orchestration.
- ▶ Collaborated with research teams to productionize ML/VLM workflows — built automated speech-to-text transcription pipelines with GPU-heavy workload serving, transforming R&D prototypes into production-scale document intelligence and pattern recognition systems.
- ▶ Designed a cloud-native, event-driven architecture using Celery and Kafka for distributed task processing and message queuing, establishing best practices for parallel processing, observability, and scalable production infrastructure.
- ▶ Built and maintained production deployments including real-time monitoring dashboards with observability best practices, surfacing KPIs, performance metrics, and enabling QC workflows with live status tracking and incident response capabilities.
- ▶ Owned infrastructure end-to-end including cloud and on-prem deployments — engineered dynamic configuration management with automated CI/CD pipelines, performance optimization, and zero-downtime updates using Docker, Kubernetes, and Helm.

AI-Driven Enterprise Data Aggregation & Lead Intelligence Platform

Scalable Python + Google Gemini Pro LLM pipeline for enterprise data extraction, enrichment, and client API delivery

- ▶ Architected and built an AI-Driven Enterprise Data Aggregation & Lead Intelligence Platform — a scalable Python-based automation system that extracts, enriches with data enrichment logic, cleans, and delivers enterprise-level company and executive data using Google Gemini Pro LLM and rule-based processing pipelines.
- ▶ Integrated Google Gemini Pro API for intelligent data parsing and enrichment, leveraging Generative AI and Prompt Engineering to classify executive profiles (CEO, VP, Director, Manager) and extract structured business attributes — job titles, employee size, revenue, and industry classification.
- ▶ Developed Python-based web scraping pipelines using Requests, BeautifulSoup, and Selenium to extract enterprise company data, LinkedIn company and individual profile URLs, and marketplace seller data from Amazon Business and Walmart, handling structured and unstructured web sources at scale.
- ▶ Implemented data cleaning, normalization, and transformation pipelines applying predefined business rules — job level classification logic, job function categorization, and deduplication workflows — ensuring data accuracy and consistency across a centralized master database.

- ▶ Managed large-scale SQL database (PostgreSQL / MySQL) updates for a high-volume master data store, optimizing query performance through indexing strategies and batch processing to support millions of enterprise records with low-latency retrieval.
- ▶ Designed and deployed RESTful APIs using Django REST Framework and FastAPI to provide client-specific filtered data access, implementing a rule-based client filter engine with JWT authentication so business clients can programmatically retrieve custom-segmented enterprise datasets.
- ▶ Formatted and serialized all processed data into structured JSON payloads for downstream consumption, enabling seamless integration with client CRM systems, analytics platforms, and Data Intelligence workflows through well-documented API contracts.
- ▶ Delivered the end-to-end solution autonomously within an agile framework, owning the full backend software stack from data extraction through API delivery, applying pragmatic problem-solving to optimize pipeline throughput and reduce data processing overhead by over 35%.

FinTech Portfolio & Trade Intelligence Platform

Full-stack Python + React.js platform for real-time portfolio tracking, trade reconciliation, and regulatory reporting on AWS

- ▶ Architected and delivered a full-stack FinTech platform for automated portfolio tracking, real-time trade reconciliation, and regulatory reporting — built with a Python/FastAPI backend, React.js frontend, and PostgreSQL data layer, deployed on AWS (EC2, RDS, Lambda).
- ▶ Developed advanced Python scripting modules for automated financial data ingestion, parsing live market feeds (REST APIs, WebSockets), normalizing transaction records, and computing real-time P&L, NAV, and risk metrics for institutional portfolios.
- ▶ Integrated LLM-powered (Generative AI) financial report summarization and anomaly detection, using Prompt Engineering and RAG pipelines to surface actionable insights from large volumes of structured trade and accounting data.
- ▶ Built responsive React.js dashboards with real-time charting (Recharts, D3.js) for portfolio visualization, risk exposure heatmaps, and executive financial summaries — ensuring sub-2-second load performance through API caching with Redis.
- ▶ Designed high-performance PostgreSQL schemas for complex financial instruments (equities, derivatives, FX), implementing audit trail tables, ACID-compliant transaction workflows, and compliance-ready data retention policies aligned with regulatory requirements.
- ▶ Managed end-to-end deployment pipeline using Docker containers, Kubernetes orchestration, and Terraform infrastructure-as-code on AWS — implementing GitHub Actions CI/CD with automated test suites achieving 95%+ code coverage.
- ▶ Integrated third-party FinTech APIs (payment gateways, banking data providers, FIX protocol adapters) with OAuth2/JWT authentication and RBAC, ensuring secure data exchange and meeting SOC 2-aligned security practices for financial data handling.

Healthcare Management & Clinical Workflow Platform

FHIR-compliant full-stack Python + React.js system for patient records, clinical automation, and AI-powered analytics on Azure

- ▶ Designed and built a full-stack Healthcare Management Platform covering patient record management, appointment scheduling, clinical workflow automation, and billing — with a Python/Django backend, React.js frontend, and deployed on Azure (App Service, Blob Storage, PostgreSQL).
- ▶ Implemented FHIR-compliant (Fast Healthcare Interoperability Resources) RESTful APIs for patient data exchange between EHR systems, enabling secure interoperability with third-party clinical platforms and ensuring HL7-standard data serialization and JSON-based record delivery.
- ▶ Developed Python scripting pipelines for automated clinical data extraction, PDF medical report parsing (PyPDF2, pdfplumber), and LLM-powered clinical note summarization using Generative AI — reducing manual data entry time by 45% for healthcare staff.
- ▶ Built an interactive React.js patient portal with real-time appointment booking, prescription tracking, and telemedicine video integration — implementing WebSocket-based notifications and responsive UI components for mobile and desktop healthcare workflows.
- ▶ Enforced HIPAA-aligned security practices across the backend architecture — end-to-end encryption (AES-256), Role-Based Access Control (RBAC) with granular clinical role permissions, audit logging, and secure data storage strategies for Protected Health Information (PHI).
- ▶ Integrated ML-based predictive analytics for patient readmission risk scoring and appointment no-show prediction using Scikit-learn models served via FastAPI inference endpoints, enabling proactive clinical intervention for high-risk patient cohorts.
- ▶ Orchestrated containerized deployment on Azure with Kubernetes, Helm charts, and Terraform — automating infrastructure provisioning, managing health monitoring with Grafana/Prometheus dashboards, and maintaining 99.95% uptime SLA for critical healthcare services.

ACHIEVEMENTS

- ▶ **[National Finalist] Smart India Hackathon 2022** — Developed an AI/ML-based Predictive Crime Analytics solution, representing Silver Oak University at the national level. Designed a forecasting model using real-time data insights to predict potential criminal activities.
- ▶ **Entrepreneurship Cell (E-Cell) – Active Member** — Actively contributed to organizing startup events, workshops, and innovation challenges, fostering entrepreneurship awareness and a culture of innovation among students.
- ▶ **E-Cell Recognition** — Recognized for active participation in E-Cell initiatives; supported entrepreneurship awareness programs and student engagement drives throughout the academic year.